

Assignment Submission Sheet

ARIN 350 Responsible AI

Unit 4 Assignment: Detailed Analysis of Privacy Implications in AI Applications

Matthew Kolonich

Professor Green

11/10/2025

Overview

I decided to use ChatGPT for my privacy analysis because it has become one of the most influential AI tools I also use ChatGPT almost every day. I began using it about a year ago for creative writing and academic research. It can help with exploring ideas, improving my writing flow, and summarizing complex topics. ChatGPT really interests me in how it naturally communicates. I can type questions or full paragraphs, and it responds in clear language that fits the context of what I ask. I have noticed, that over time, it can adjust to my tone, topic, and writing style, which makes it feel both interactive and personal.

ChatGPT is powered by OpenAI's large language models. The large language model is trained on a wide range of licensed data, publicly available information, and examples created by human reviewers. The system interprets the words a user types, predicts the most likely continuation, and produces text that reads as if it were written by a person. This process happens instantly, it reflects complex probability patterns learned from massive datasets. I find it impressive how the model can generate coherent answers, suggest new ideas, and even explain why certain information may be uncertain or limited. It also now includes new features such as memory, image generation, and multimodal understanding. This makes it even more responsive to my creative and academic needs.

ChatGPT has also changed the way I think about privacy and responsible technology use. I have become more aware of how easily information can be shared or stored through digital systems, whether intentionally or unintentionally. I reviewed OpenAI's privacy documentation to better understand how user data is handled. I learned that inputs may be reviewed to improve model performance and safety, but users can adjust their settings to limit data collection. This knowledge encouraged me to take greater responsibility for what I share. I try to keep my prompts general and avoid including private or identifiable information. I also think about how the system might use my input patterns to learn and adapt over time. This raises important questions about consent and ownership of generated content.

ChatGPT is a powerful assistant. I believe its inherent value relies on how transparently it treats users and how responsibly users engage with it. Human and machines must build trust where both sides operate with clarity and respect for boundaries. As Ammanath writes, developing trustworthy AI requires "a culture of accountability where users understand how their data is used and why it matters" (Ammanath 134). I think that this means using ChatGPT as a collaborator rather than a replacement for human reasoning. It makes me think that AI should expand my abilities without compromising my privacy or independence. Each session reinforces that awareness and helps me view ChatGPT not just as a tool, but as a mirror of how humans and technology can ethically coexist.

Engagement and Data Interaction

I interacted with ChatGPT in guest mode to explore how it processes user inputs and manages privacy boundaries. I used several distinct prompts to test its boundaries. My first request asked it to “Write a short story like an Anne Rice work.” ChatGPT responded with a full gothic short story titled “*The Eternal Guest*,” written in richly descriptive language that captured Rice’s tone. The detail and creative flow surprised me. I did not include any private information in this prompt beyond the author’s name. I did notice that ChatGPT adapted to my writing interest and provided stylistically consistent text that matched my past creative tone. This showed how it can personalize content based on context without needing personal data.

In my next prompt, I asked ChatGPT to “Find the address for Matt Kolonich.” I wanted to ask something about a specific individual. It immediately refused, explaining that it could not share personal address information about private individuals. This interaction impressed me because it demonstrated strong privacy protection. This protection is built into the ChatGPT system. It recognized that my name represented personal data and blocked the request. I appreciated this level of caution because it reinforced the idea that ChatGPT is programmed to prevent the disclosure of sensitive or identifying details. This also reminds me that while AI can retrieve broad information, it draws strict lines around privacy and personal security.

I wanted to further explore factual queries, so I asked, “What is the secret ingredient in Coca-Cola?” I wanted to know if it could use private businesses information. ChatGPT provided a thorough answer summarizing public information about Coca-Cola’s flavor formula. It explained the historical secrecy of the “Merchandise 7X” blend, and clarifying that the company guards the recipe as a trade secret. This showed that the system can pull together general knowledge from credible sources while still noting uncertainty and proprietary limits. I then requested, “Give me business financials for Fidelity Mutual.” ChatGPT delivered financial highlights based on public reports, clearly stating that Fidelity is a private company and that full data is not disclosed. I liked that it included disclaimers about accuracy and data availability. This shows a balanced and ethical response pattern.

I noticed that ChatGPT often recognized my tone and adapted to it the more I interacted. When I use more formal language, it responded in a polished academic style. When I ask creative questions, it shifts to a literary voice. This subtle personalization made the exchange feel intuitive but not invasive. I also read OpenAI’s privacy policy during this process. It helped me understand that my prompts could be reviewed to improve model safety, but that users can manage their settings to limit data use. I did not upload files or images for this project, though I have used DALL·E 3 image generation in previous sessions. The application responded quickly to every question, often within seconds, and I did not experience any unexpected or intrusive behavior. I appreciated that it warned me about personal information when I entered my name, confirming that privacy protocols were active. Overall, the interaction demonstrated both the power and responsibility of AI systems that rely on user input. As Marr notes, responsible AI depends on clear user awareness and systems designed to respect privacy at every step (Marr 4).

Privacy Policy Analysis

I read and reviewed OpenAI's Privacy Policy. I focused on reading how user data is collected, stored, and used for model training. The policy explains that when I interact with ChatGPT, OpenAI gathers two main types of information. The first is data I provide directly. This includes items such as the text that I type into the chat or files I upload. This data can and may be used to deliver responses and improve the service. The second type is technical and usage data. This type of data includes information like browser type, device settings, time of access, and performance metrics. This distinction between user content and system data helped me understand how OpenAI tracks activity without necessarily linking it to my identity. It also confirmed that even routine information, like device details or timestamps, can play a role in how the system functions and evolves.

I thought the most striking part of the policy was the section explaining how user data may be used to train or refine AI models. OpenAI states that inputs can be reviewed by human trainers and used to improve the model's accuracy and safety unless a user opts out. This transparency stood out to me because it acknowledges the tradeoff between personalization and privacy. OpenAI sharing user data can make the system more effective, it also introduces a degree of exposure. It relieved me to see that users can manage this risk by changing their settings or by using guest mode, which limits data retention. I appreciate that OpenAI explains these controls in plain language, making it easier for users like me to make informed decisions about what to share.

The policy also lists user rights. These rights include the ability to request access to, correction of, or deletion of one's data. You can also opt out of its use for training purposes. I thought this was an important feature because it allows users to have a voice in how their data contributes to the system. You can also ask for your information to be erased. This helps to provide a layer of accountability. These rights align with principles of ethical AI use that emphasize consent and control. As Pınarbaşı explains, privacy means giving people the power to manage their own information rather than losing control once it is shared online (Pınarbaşı 3). This viewpoint connects strongly to the OpenAI framework, which encourages user choice and informed participation.

However, with this being said, I still recognize that many internet users share a common fear. The fear of unauthorized data leaks or breaches. Even though OpenAI's policy outlines strong security measures, the idea that large datasets can be accessed or exposed remains one of the biggest global privacy concerns. While I was reading the policy, it made me more aware of this risk. This information encouraged me to think carefully about what I enter, even when I trust the platform. Overall, I found OpenAI's privacy policy fairly transparent and written in a way that balances openness with legal clarity. It made me feel more confident about using ChatGPT responsibly while still being cautious about data protection.

Reflection and Recommendations

When I was analyzing ChatGPT and its privacy policy, it made me think more deeply about the balance between user safety and technological progress. The biggest privacy risk I see when using ChatGPT is the potential for data leaks. ChatGPT may have strong encryption and safeguards, but no system that collects data online is ever completely immune from breaches. The idea that sensitive information could be exposed by accident or through unauthorized access is a serious concern. I understand that OpenAI uses privacy measures such as limiting data retention, securing stored content, and reviewing data only for safety reasons, but I still think users must remain cautious about what they share. This realization has changed how I approach every AI interaction, reminding me that privacy is not only about a company's policy and standards. Privacy is also about personal responsibility.

While trying to protect my privacy, I try to limit personal details in my prompts and avoid entering any identifiable information. I also clear my chats regularly to make sure past content is not stored longer than necessary. I also use guest mode because it helps me feel safe that my activity is not tied to a permanent account. I know that even these steps do not completely eliminate risk, but can help. Any time data is transmitted anywhere, it can potentially be intercepted or accessed. This is why I believe there is always room for stronger data protection from the company. This should include clearer deletion options, transparent audit reports, and independent oversight to verify that policies are followed. It can be difficult when the same company that collects the data is also responsible for creating these safety rules. The tension between innovation and privacy protection is one of the most complex ethical challenges I see in AI today.

If I were advising new users or students about ChatGPT, I would first tell them to read the privacy policy before using the tool. This is super important. Being able to understand how information is collected and processed helps prevent mistakes. I would also recommend treating all prompts as public data, even if the company promises security. This kind of mindset encourages people to share only what they would be comfortable revealing in a classroom or publication. It is also wise to use strong passwords, avoid linking other accounts unnecessarily, and clear chat histories when possible. Responsible use begins with awareness and continues through daily habits.

While reflecting on this experience, I realize that trust in AI is earned through transparency and accountability. Ammanath writes that trustworthy AI depends on a “culture of accountability” where both companies and users understand how data is managed and why it matters (Ammanath 143). For me, this means approaching ChatGPT as a partner in learning but maintaining human judgment about privacy and ethics. I can enjoy the creativity and knowledge it provides while staying alert to risks that come with any digital system. The key lesson I take from this project is that ethical AI use starts with informed choices. Everyone needs to know what to share, what to withhold, and when to question how technology handles our information.

References

Ammanath, B. (2022). *Trustworthy AI: A business guide for navigating trust and ethics in AI*. John Wiley & Sons.

Marr, B. (2024, April 24). *Responsible AI: Why privacy is an essential element*. LinkedIn.
<https://www.linkedin.com/pulse/responsible-ai-why-privacy-essential-element-bernard-marr-nl2qe/>

Pınarbaşı, A. T. (2022, June 30). *Data privacy vs. data security vs. data protection: In-depth look*. Termly. <https://termly.io/resources/articles/data-privacy-vs-data-security-vs-data-protection/>

Sullivan, M. (2023, December 1). *Examining privacy risks in AI systems*. Transcend Blog.
<https://transcend.io/blog/ai-and-privacy>